

In [1]:

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

In [2]:

```
df2=pd.read_csv(r'C:\Users\user\Downloads\C2_test.gender_submission.csv')
df2
fin2=pd.read_csv(r'C:\Users\user\Downloads\C2_train.gender_submission.csv')
fin2
```

Out[2]:

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare
0	1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171	7.250
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...)	female	38.0	1	0	PC 17599	71.283
2	3	1	3	Heikkinen, Miss. Laina	female	26.0	0	0	STON/O2. 3101282	7.925
3	4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	0	113803	53.100
4	5	0	3	Allen, Mr. William Henry	male	35.0	0	0	373450	8.050
...
886	887	0	2	Montvila, Rev. Juozas	male	27.0	0	0	211536	13.000
887	888	1	1	Graham, Miss. Margaret Edith	female	19.0	0	0	112053	30.000
888	889	0	3	Johnston, Miss. Catherine Helen "Carrie"	female	NaN	1	2	W./C. 6607	23.450
889	890	1	1	Behr, Mr. Karl Howell	male	26.0	0	0	111369	30.000
890	891	0	3	Dooley, Mr. Patrick	male	32.0	0	0	370376	7.750

891 rows × 12 columns



In [3]:

```
df0=df2.head(10)
df=df0.fillna("30.0")
df
fin0=fin2.head(10)
fin=fin0.fillna("35.0")
fin
```

Out[3]:

PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare
0	1	0	3Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171	7.2500
1	2	1	1Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	0	PC 17599	71.2833
2	3	1	3Heikkinen, Miss. Laina	female	26.0	0	0	STON/O2. 3101282	7.9250
3	4	1	1Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	0	113803	53.1000
4	5	0	3Allen, Mr. William Henry	male	35.0	0	0	373450	8.0500
5	6	0	3Moran, Mr. James	male	35.0	0	0	330877	8.4583
6	7	0	1McCarthy, Mr. Timothy J	male	54.0	0	0	17463	51.8625
7	8	0	3Palsson, Master. Gosta Leonard	male	2.0	3	1	349909	21.0750
8	9	1	3Johnson, Mrs. Oscar W (Elisabeth Vilhelmina Berg)	female	27.0	0	2	347742	11.1333
9	10	1	2Nasser, Mrs. Nicholas (Adele Achem)	female	14.0	1	0	237736	30.0708



In [4]:

```
df.info()  
fin.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 10 entries, 0 to 9  
Data columns (total 11 columns):  
#   Column          Non-Null Count  Dtype  
---  -  
0   PassengerId      10 non-null     int64  
1   Pclass           10 non-null     int64  
2   Name             10 non-null     object  
3   Sex              10 non-null     object  
4   Age              10 non-null     float64  
5   SibSp            10 non-null     int64  
6   Parch            10 non-null     int64  
7   Ticket           10 non-null     object  
8   Fare             10 non-null     float64  
9   Cabin            10 non-null     object  
10  Embarked         10 non-null     object  
dtypes: float64(2), int64(4), object(5)  
memory usage: 1008.0+ bytes  
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 10 entries, 0 to 9  
Data columns (total 12 columns):  
#   Column          Non-Null Count  Dtype  
---  -  
0   PassengerId      10 non-null     int64  
1   Survived         10 non-null     int64  
2   Pclass           10 non-null     int64  
3   Name             10 non-null     object  
4   Sex              10 non-null     object  
5   Age              10 non-null     object  
6   SibSp            10 non-null     int64  
7   Parch            10 non-null     int64  
8   Ticket           10 non-null     object  
9   Fare             10 non-null     float64  
10  Cabin            10 non-null     object  
11  Embarked         10 non-null     object  
dtypes: float64(1), int64(5), object(6)  
memory usage: 1.1+ KB
```

In [5]:

```
df.describe()  
fin.describe()
```

Out[5]:

	PassengerId	Survived	Pclass	SibSp	Parch	Fare
count	10.00000	10.000000	10.000000	10.000000	10.000000	10.000000
mean	5.50000	0.500000	2.300000	0.700000	0.300000	27.020820
std	3.02765	0.527046	0.948683	0.948683	0.674949	23.601938
min	1.00000	0.000000	1.000000	0.000000	0.000000	7.250000
25%	3.25000	0.000000	1.250000	0.000000	0.000000	8.152075
50%	5.50000	0.500000	3.000000	0.500000	0.000000	16.104150
75%	7.75000	1.000000	3.000000	1.000000	0.000000	46.414575
max	10.00000	1.000000	3.000000	3.000000	2.000000	71.283300

In [6]:

```
df.columns  
fin.columns
```

Out[6]:


```
Index(['PassengerId', 'Survived', 'Pclass', 'Name', 'Sex', 'Age', 'SibSp',  
      'Parch', 'Ticket', 'Fare', 'Cabin', 'Embarked'],  
      dtype='object')
```

In [7]:

```
sns.pairplot(df)
sns.pairplot(fin)
```

Out[7]:

<seaborn.axisgrid.PairGrid at 0x2918b32d340>





```
sns.distplot(df['Age'])
sns.distplot(fin['Age'])
```

C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).


warnings.warn(msg, FutureWarning)

C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

warnings.warn(msg, FutureWarning)

Out[8]:

<AxesSubplot:xlabel='Age', ylabel='Density'>




In [9]:

```
sns.histplot(df["Age"],color='red')
sns.histplot(fin["Parch"],color='blue')
```

Out[9]:

<AxesSubplot:xlabel='Age', ylabel='Count'>



In [10]:


```
df1=df[['PassengerId', 'Pclass', 'Name', 'Sex', 'Age', 'SibSp', 'Parch',
        'Ticket', 'Fare', 'Cabin', 'Embarked']]
fin1=fin[['PassengerId', 'Pclass', 'Name', 'Sex', 'Age', 'SibSp', 'Parch',
        'Ticket', 'Fare', 'Cabin', 'Embarked']]
```

In [11]:

```
sns.heatmap(df1.corr())
sns.heatmap(fin1.corr())
```

Out[11]:

<AxesSubplot:>



In [12]:

```
x=df1[['PassengerId', 'Pclass', 'Age', 'SibSp', 'Parch', 'Fare']]
y=fin1[['PassengerId', 'Pclass', 'Age', 'SibSp', 'Parch', 'Fare']]
```

In [13]:

```
from sklearn.model_selection import train_test_split
```

In [14]:

```
x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.3)
```

In [15]:

```
from sklearn.linear_model import LinearRegression

lr=LinearRegression()
lr.fit(x_train,y_train)#ValueError: Input contains NaN, infinity or a value too large for
```

Out[15]:

```
LinearRegression()
```

In [16]:

```
print(lr.intercept_)
```

```
[ -891.          332.87835864 -4954.83029001   138.61697769
    0.          -6934.40841297]
```

In [17]:

```
coef= pd.DataFrame(lr.coef_)
coef
```

Out[17]:

	0	1	2	3	4	5
0	1.000000	-4.866403e-16	6.924589e-16	5.515729e-16	0.0	3.279364e-15
1	-0.354184	-3.868372e+00	-9.726923e-02	2.681135e-01	0.0	5.464786e-02
2	5.532481	1.657792e+01	5.239007e-01	1.407403e-01	0.0	-3.184044e+00
3	-0.161799	1.715021e+00	2.249348e-02	-2.257961e-01	0.0	1.245994e-01
4	0.000000	0.000000e+00	0.000000e+00	0.000000e+00	0.0	0.000000e+00
5	7.516158	7.225946e+01	1.859503e+00	1.470512e+01	0.0	-3.811048e+00

In [18]:

```
print(lr.score(x_test,y_test))
```


```
-22.52209907188785
```

In [19]:

```
prediction = lr.predict(x_test)
plt.scatter(y_test,prediction)
```

Out[19]:

<matplotlib.collections.PathCollection at 0x2918bbb62b0>



In [20]:

```
lr.score(x_test,y_test)
```

Out[20]:

-22.52209907188785

In [21]:

```
lr.score(x_train,y_train)
```

Out[21]:

0.8916852718919287

In [22]:

```
from sklearn.linear_model import Ridge,Lasso
```

In [23]:

```
rr=Ridge(alpha=10)
rr.fit(x_train,y_train)
```

Out[23]:

Ridge(alpha=10)

In [24]:

```
rr.score(x_test,y_test)
```

Out[24]:

-3.383452576892554

In [25]:

```
la=Lasso(alpha=10)
la.fit(x_train,y_train)
```

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model_coordinate_descent.py:530: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 0.0, tolerance: 0.0

```
model = cd_fast.enet_coordinate_descent(
```

Out[25]:

```
Lasso(alpha=10)
```

In [26]:

```
la.score(x_test,y_test)
```

Out[26]:

```
-2.207084556972902
```

Elastic Net

In [27]:

```
from sklearn.linear_model import ElasticNet
en = ElasticNet()
en.fit(x_train,y_train)
```

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model_coordinate_descent.py:530: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 0.0, tolerance: 0.0

```
model = cd_fast.enet_coordinate_descent(
```

Out[27]:

```
ElasticNet()
```

In [28]:

```
print(en.coef_)
```

```
[[ 7.38992951e-01 -0.00000000e+00 -1.95680928e-02  0.00000000e+00
  0.00000000e+00  6.77545964e-02]
 [-0.00000000e+00 -0.00000000e+00  2.08752798e-03 -0.00000000e+00
  0.00000000e+00  1.61910530e-03]
 [ 3.68143731e+00  7.97420625e-02  6.51140915e-02  5.10797984e-01
  0.00000000e+00 -2.78083821e+00]
 [-0.00000000e+00  0.00000000e+00 -1.77696945e-05  0.00000000e+00
  0.00000000e+00  7.83565915e-03]
 [ 0.00000000e+00  0.00000000e+00  0.00000000e+00  0.00000000e+00
  0.00000000e+00  0.00000000e+00]
 [ 3.42014506e+00  5.01721588e+00  2.76294345e-01  9.34754689e+00
  0.00000000e+00 -2.71379332e+00]]
```

In [29]:

```
print(en.intercept_)
```

```
[-6.57271969e+02  1.91261092e+00 -3.23868308e+03  4.88986989e-01  
 0.00000000e+00 -3.02949935e+03]
```

In [30]:

```
prediction=en.predict(x_test)  
print(prediction)
```

```
[[ 5.26775162  1.97843129 27.89774242  0.58487672  0.          32.08255313]  
 [ 8.5388068  2.0138407 -7.35198985  0.71575909  0.          -6.92332105]  
 [ 7.95927272  1.96189126 55.91855122  0.54531268  0.          49.0375898  
 3]]
```

In [31]:

```
print(en.score(x_test,y_test))
```

```
-5.080616467467732
```

Logistic Regression

In [32]:

```
from sklearn.linear_model import LogisticRegression
```

In [33]:

```
feature_matrix=df1[['PassengerId', 'Pclass', 'Age', 'SibSp', 'Parch','Fare']]  
target_vector = df1[['Embarked']]
```

In [34]:

```
feature_matrix.shape
```

Out[34]:

```
(10, 6)
```

In [35]:

```
target_vector.shape
```

Out[35]:

```
(10, 1)
```

In [36]:

```
from sklearn.preprocessing import StandardScaler
```

In [37]:

```
fs = StandardScaler().fit_transform(feature_matrix)
```

In [38]:

```
logr = LogisticRegression()  
logr.fit(fs,target_vector)
```

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\utils\validation.py:63:
DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().
return f(*args, **kwargs)

Out[38]:

```
LogisticRegression()
```

In [41]:

```
observation=df1[['PassengerId', 'Pclass', 'Age', 'SibSp', 'Parch','Fare']]
```

In [42]:

```
prediction = logr.predict(observation)  
print(prediction)
```

```
['C' 'C' 'C' 'C' 'C' 'C' 'C' 'C' 'C' 'C']
```

In [43]:

```
logr.classes_
```

Out[43]:

```
array(['C', 'Q', 'S'], dtype=object)
```

In [44]:

```
logr.predict_proba(observation)[0][0]
```

Out[44]:

```
1.0
```

In [45]:

```
logr.predict_proba(observation)[0][1]
```

Out[45]:

```
0.0
```

Evaluation Metrics

In [46]:

```
from sklearn import metrics
```

In [47]:

```
print("Mean Absolute Error:", metrics.mean_absolute_error(y_test, prediction))
```

```
-----
-
ValueError                                Traceback (most recent call last)
<ipython-input-47-4b3ebc69c71d> in <module>
----> 1 print("Mean Absolute Error:", metrics.mean_absolute_error(y_test, pr
ediction))

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\utils\validation.py in
inner_f(*args, **kwargs)
    61         extra_args = len(args) - len(all_args)
    62         if extra_args <= 0:
--> 63             return f(*args, **kwargs)
    64
    65         # extra_args > 0

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\metrics\_regression.py
in mean_absolute_error(y_true, y_pred, sample_weight, multioutput)
    180     0.85...
    181     """
--> 182     y_type, y_true, y_pred, multioutput = _check_reg_targets(
    183         y_true, y_pred, multioutput)
    184     check_consistent_length(y_true, y_pred, sample_weight)

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\metrics\_regression.py
in _check_reg_targets(y_true, y_pred, multioutput, dtype)
    86         the dtype argument passed to check_array.
    87         """
--> 88     check_consistent_length(y_true, y_pred)
    89     y_true = check_array(y_true, ensure_2d=False, dtype=dtype)
    90     y_pred = check_array(y_pred, ensure_2d=False, dtype=dtype)

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\utils\validation.py in
check_consistent_length(*arrays)
    260     uniques = np.unique(lengths)
    261     if len(uniques) > 1:
--> 262         raise ValueError("Found input variables with inconsistent
numbers of"
    263                               " samples: %r" % [int(l) for l in length
s])
    264

ValueError: Found input variables with inconsistent numbers of samples:
[3, 10]
```

In [48]:

```
print("Mean Squared Error:", metrics.mean_squared_error(y_test, prediction))
```

```
-----
-
ValueError                                Traceback (most recent call last)
<ipython-input-48-a8cbb0b7f78a> in <module>
----> 1 print("Mean Squared Error:", metrics.mean_squared_error(y_test, prediction))

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\utils\validation.py in
inner_f(*args, **kwargs)
    61         extra_args = len(args) - len(all_args)
    62         if extra_args <= 0:
--> 63             return f(*args, **kwargs)
    64
    65         # extra_args > 0

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\metrics\_regression.py
in mean_squared_error(y_true, y_pred, sample_weight, multioutput, squared)
    333     0.825...
    334     """
--> 335     y_type, y_true, y_pred, multioutput = _check_reg_targets(
    336         y_true, y_pred, multioutput)
    337     check_consistent_length(y_true, y_pred, sample_weight)

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\metrics\_regression.py
in _check_reg_targets(y_true, y_pred, multioutput, dtype)
    86         the dtype argument passed to check_array.
    87         """
--> 88     check_consistent_length(y_true, y_pred)
    89     y_true = check_array(y_true, ensure_2d=False, dtype=dtype)
    90     y_pred = check_array(y_pred, ensure_2d=False, dtype=dtype)

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\utils\validation.py in
check_consistent_length(*arrays)
    260     uniques = np.unique(lengths)
    261     if len(uniques) > 1:
--> 262         raise ValueError("Found input variables with inconsistent
numbers of"
    263                             " samples: %r" % [int(l) for l in length
s])
    264

ValueError: Found input variables with inconsistent numbers of samples:
[3, 10]
```

In [49]:

```
print("Root Mean Squared Error:", np.sqrt(metrics.mean_squared_error(y_test, prediction)))
```

```
-----
-
ValueError                                Traceback (most recent call last)
<ipython-input-49-c32fd676945a> in <module>
----> 1 print("Root Mean Squared Error:", np.sqrt(metrics.mean_squared_error(y_test, prediction)))
```

```
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\utils\validation.py in inner_f(*args, **kwargs)
    61         extra_args = len(args) - len(all_args)
    62         if extra_args <= 0:
--> 63             return f(*args, **kwargs)
    64
    65         # extra_args > 0
```

```
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\metrics\_regression.py in mean_squared_error(y_true, y_pred, sample_weight, multioutput, squared)
    333     0.825...
    334     """
--> 335     y_type, y_true, y_pred, multioutput = _check_reg_targets(
    336         y_true, y_pred, multioutput)
    337     check_consistent_length(y_true, y_pred, sample_weight)
```

```
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\metrics\_regression.py in _check_reg_targets(y_true, y_pred, multioutput, dtype)
    86         the dtype argument passed to check_array.
    87         """
--> 88     check_consistent_length(y_true, y_pred)
    89     y_true = check_array(y_true, ensure_2d=False, dtype=dtype)
    90     y_pred = check_array(y_pred, ensure_2d=False, dtype=dtype)
```

```
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\metrics\_regression.py in _check_reg_targets(y_true, y_pred, multioutput, dtype)
    86         the dtype argument passed to check_array.
    87         """
--> 88     check_consistent_length(y_true, y_pred)
    89     y_true = check_array(y_true, ensure_2d=False, dtype=dtype)
    90     y_pred = check_array(y_pred, ensure_2d=False, dtype=dtype)
```

```
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\metrics\_regression.py in _check_reg_targets(y_true, y_pred, multioutput, dtype)
    86         the dtype argument passed to check_array.
    87         """
--> 88     check_consistent_length(y_true, y_pred)
    89     y_true = check_array(y_true, ensure_2d=False, dtype=dtype)
    90     y_pred = check_array(y_pred, ensure_2d=False, dtype=dtype)
```

```
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\metrics\_regression.py in _check_reg_targets(y_true, y_pred, multioutput, dtype)
    86         the dtype argument passed to check_array.
    87         """
--> 88     check_consistent_length(y_true, y_pred)
    89     y_true = check_array(y_true, ensure_2d=False, dtype=dtype)
    90     y_pred = check_array(y_pred, ensure_2d=False, dtype=dtype)
```

```
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\metrics\_regression.py in _check_reg_targets(y_true, y_pred, multioutput, dtype)
    86         the dtype argument passed to check_array.
    87         """
--> 88     check_consistent_length(y_true, y_pred)
    89     y_true = check_array(y_true, ensure_2d=False, dtype=dtype)
    90     y_pred = check_array(y_pred, ensure_2d=False, dtype=dtype)
```

```
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\utils\validation.py in check_consistent_length(*arrays)
    260     uniques = np.unique(lengths)
    261     if len(uniques) > 1:
--> 262         raise ValueError("Found input variables with inconsistent
numbers of"
    263                             " samples: %r" % [int(l) for l in length
s])
    264
```

```
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\utils\validation.py in check_consistent_length(*arrays)
    260     uniques = np.unique(lengths)
    261     if len(uniques) > 1:
--> 262         raise ValueError("Found input variables with inconsistent
numbers of"
    263                             " samples: %r" % [int(l) for l in length
s])
    264
```

```
ValueError: Found input variables with inconsistent numbers of samples:
[3, 10]
```

Model Saving

In [50]:

```
import pickle
```


In [51]:

```
filename="prediction"  
pickle.dump(lr,open(filename,'wb'))
```